

IoT Lab 3

Multi-hop Collection



Overview

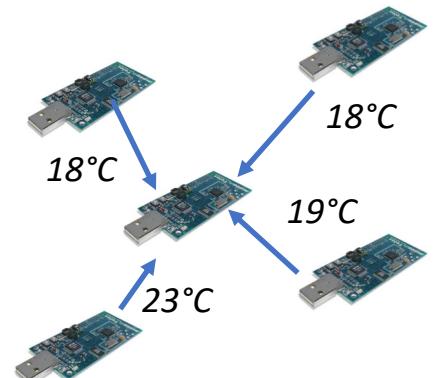
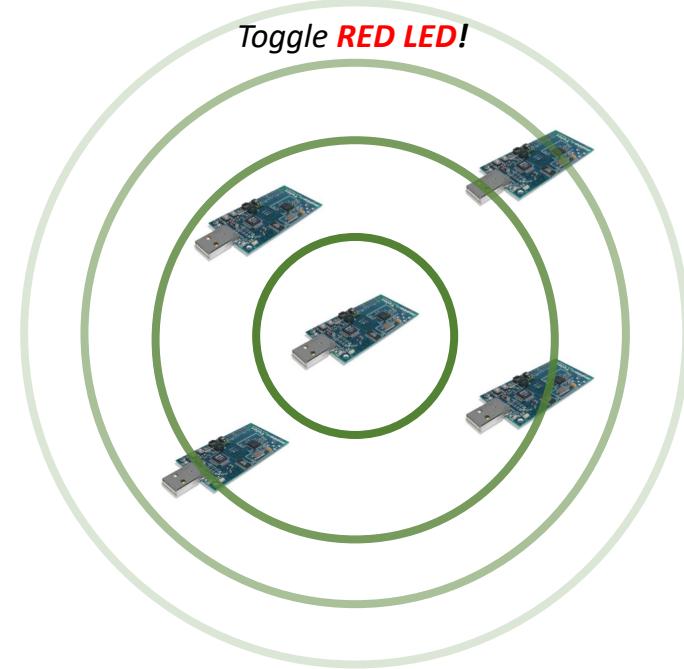
- Basic concept
- Setup
- Tasks
- Hints
- Your deliverables

Basic Concept



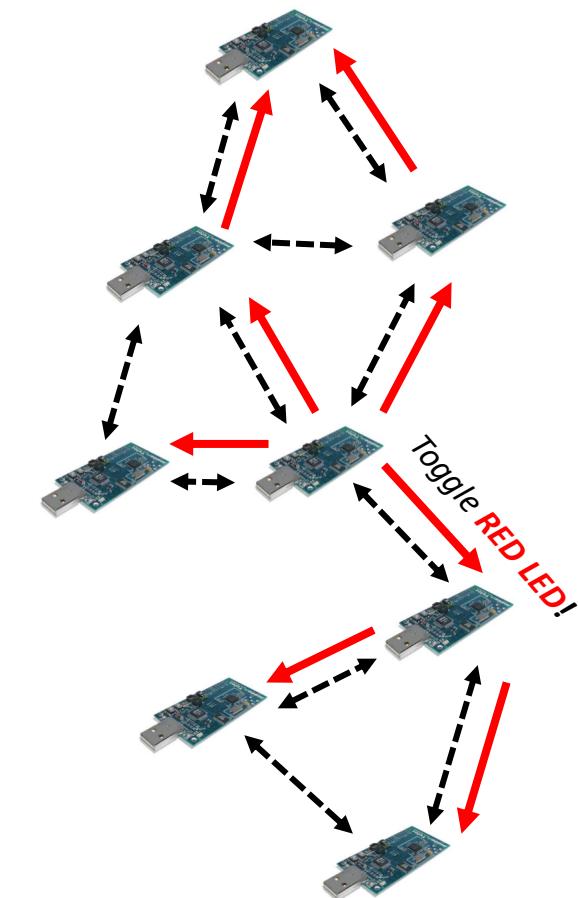
Summary of Lab 1

- In the first lab, you saw two fundamental communication models:
 - One-to-all / Dissemination
 - delivering commands
 - Network-wide configuration/code update
 - All-to-one / Collection
 - collecting sensor data to a gateway
 - Alarm detection
- However, we limited ourselves to one-hop
 - Everybody can contact everybody else
 - Many deployments are multi-hop!
 - Ex: Avalanche detection, wildlife monitoring, large factories



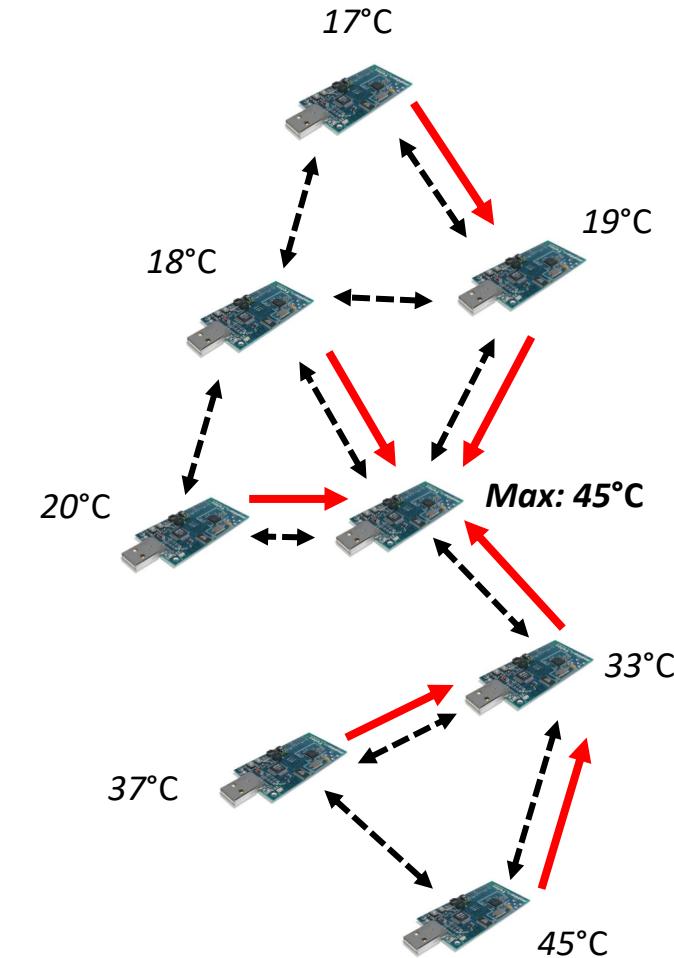
Summary of Lab 2

- Lab 2 introduced multi-hop topologies and you have implemented a multi-hop dissemination solution
 - Multi-hop: a node can not reach all other nodes of the network
- Multi-hop dissemination is used to:
 - Disseminate commands, new software, control packets
 - Build routes in AODV
 - Synchronize clocks, etc.
- One simple solution?
 - Flooding using broadcasts!



Lab 3: basic concept

- This lab will focus on **Multi-hop collection**
 - Again, a **central entity** can not reach all nodes and must rely on relays
- What is the goal?
 - **Collect the maximum** temperature recorded by a network of sensors
 - Be as efficient as possible
- Again, we will evaluate our protocol

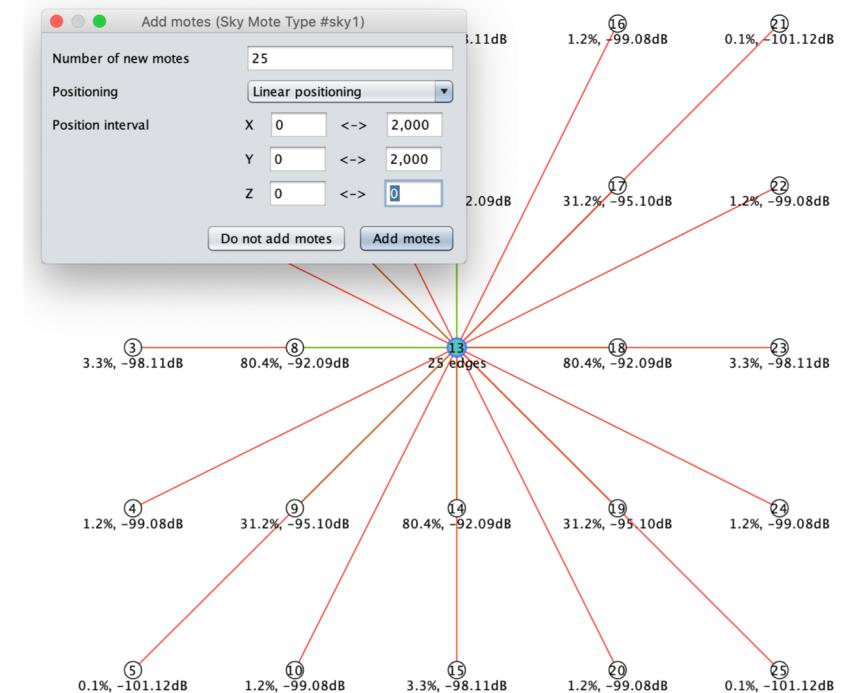


Setup



Initial Setup

- Create a new Cooja simulation
 - Multi-path Ray-tracer Medium (MRM)
 - Mote startup delay: 0 ms
 - 25 Sky motes
 - Linear positioning
 - X: 0-2000
 - Y: 0-2000
- By clicking on nodes, you should see that the probability of receiving messages depends on the distance (+ message collision)



Tasks

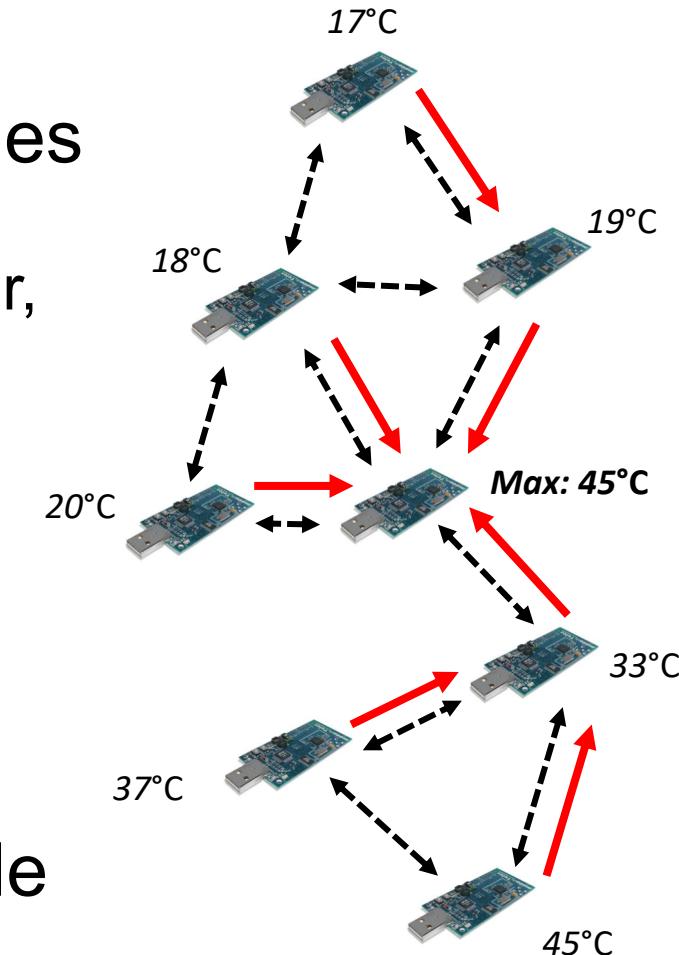
What you must do



Part 1 – Temperature collection

25 nodes: one central node (ID 13), 24 "slave" nodes (ID 1-12 and 14-25)

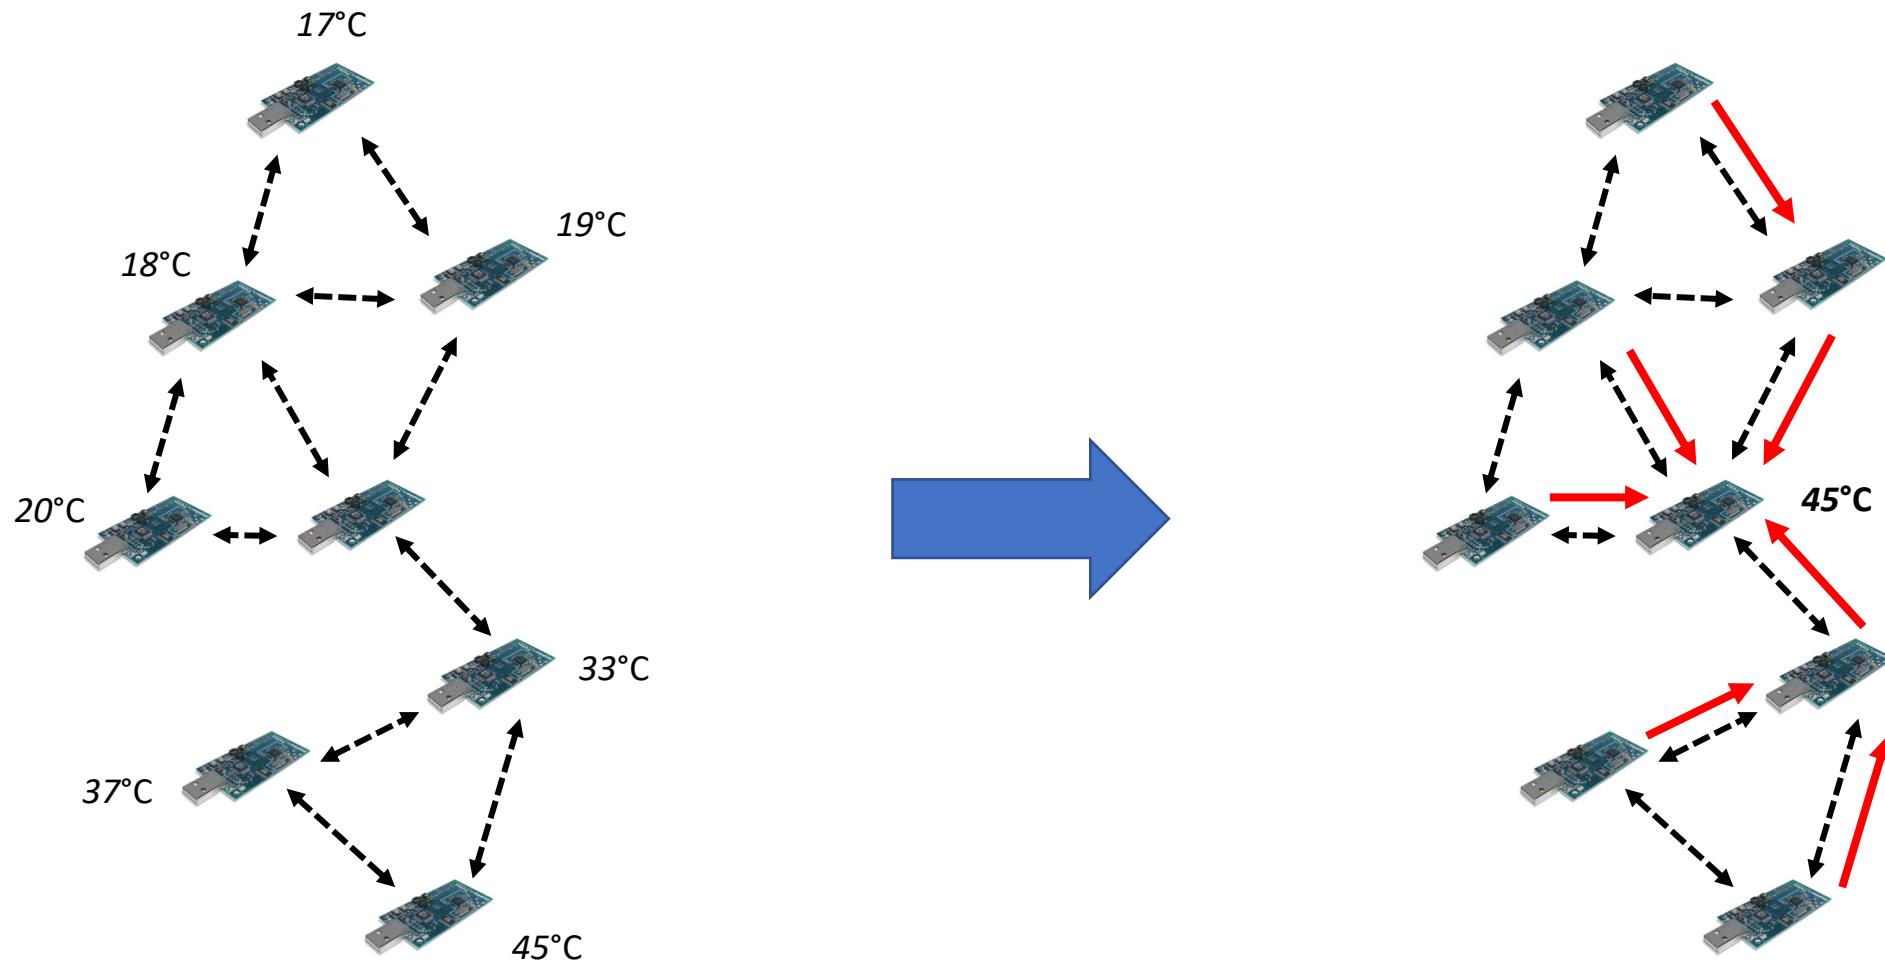
- If you followed the setup, ID 13 should be at the center, and have a 80% probability to communicate with the nearest nodes



Every 2 minutes:

- Slave nodes generate a **random** temperature between -100 and 100 degrees
- The temperature is propagated to the central node
- The central node displays the **maximum** temperature received
 - Unlike Lab 1, the central node does not have to print all received values

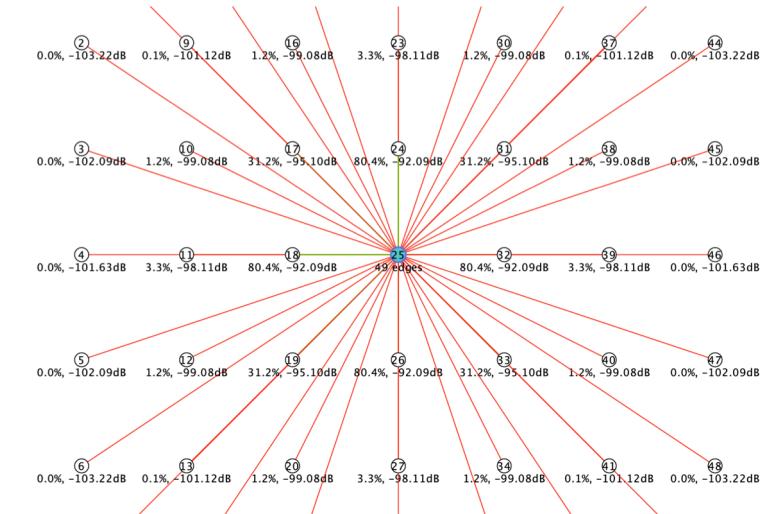
Example



Part 2 – Evaluation (1)

You will evaluate the reliability and latency of your solution, for networks with increasing diameter (= number of hops)

- Network sizes:
 - 9 nodes (linear, max: 1200, central node: 5)
 - 25 nodes (linear, max: 2000, central node: 13)
 - 49 nodes (linear, max: 2800, central node: 25)
- What to evaluate?
 - Latency: for one command, time until the last node received the command
 - Reliability: for one command, how many nodes received it?



Part 2 – Evaluation (2)

- What is a (good) evaluation?
 - A figure/plot, not just numbers thrown around
 - Should be representative of a typical execution
 - Should present the statistical aspect of running the protocol more than once (what is the mean, standard deviation or quartiles ?)
 - You should run multiple rounds and present us your methodology!!!
- How to collect the data?
 - Use printf() or LOG_INFO()
 - Cooja's mote output can be exported as .txt file (Mote output: File->save to file)
 - The output file will contain cooja's global time

Hints



Some hints

- How to solve this lab?
 - There are many approaches possible for part 1, it is up to you how to design your preferred solution
 - Some common ways include:
 - Build a routing tree, and then use unicasts
 - Use flooding for each value
 - Use in-network processing (i.e., aggregate data along the way)
 - We give you additional reading materials that you can get inspiration from
- How to avoid the initial random initialization for part 2?
 - When creating the simulation, set the *mote startup delay* to 0 ms

Your deliverables



You must submit

- A **video** presenting your solutions:
 - 5 to 8 min (we remove points if the video is longer than 8 min!)
 - The video must contain:
 - A demo of your solution in cooja
 - An explanation of:
 - your design choices
 - your code
 - possible corner cases (what could break your solution)?
 - **Your evaluation**
 - Plots/figures (with axes information, how many runs were used per datapoints, etc.)
 - explanation of the main findings
 - **Your video does not have to be perfect, don't spend too much time editing, a few pauses and hesitations are perfectly normal**
- **Your code**
 - Any file you might have modified
- **Your figures/plots**

How to submit?

- Upload your source code, plots and video as an archive (zip, rar) on **iLearn**
 - iLearn limits the file size to ~20Mb, if your file is too large, upload the source code and plots only on iLearn, and the video on a dropbox link given in iLearn
 - Please note that the dropbox repository is shared with everybody, so you must limit your video size to **50Mb!!!**
 - **Additionaly, you can now submit the VIDEO as a dropbox link (using your own account, no size limit) or on Youtube.**
- Like the prelab, you will work as a group (2 students), if you submit alone you need permission from us

Good luck!

Deadline: Monday, 25th May



To go further

- *O. Gnawali, R. Fonseca, et al.* "Collection tree protocol". 2009
- RFC 6550, RPL: IPv6 Routing Protocol for Low-Power and Lossy Networks, 2012
- *F. Ferrari, M. Zimmerling, et al.* "Low-power wireless bus". 2012
- *O. Landsiedel, F. Ferrari, et al.* "Chaos: versatile and efficient all-to-all data sharing and in-network processing at scale". 2013